

The Overcorrection Paradox

WHAT THEY DO

64.3%

of the time, models revise
output already rated sufficient

yet

WHAT THEY KNOW

93.3%

of the time, models prefer
Turn 1 over Turn 5 in blind test

Models know revision makes output worse, but revise anyway.