

Study 3: Revision Yield Experiment

720 trials

6 models

5 domains

40 tasks

Task Prompt
e.g., 'Write a debounce function'

T1: 4.27
above thresh.

Working Model Generates Output
(Turn N, temperature 1.0)

T5: 3.04
below thresh.

Attrition

T1: 720
T2: 491 (68%)
T3: 362 (50%)
T4: 304 (42%)
T5: 253 (35%)

Blind Evaluator Scores Output
1-6 quality scale | separate Claude instance

BLIND
Threshold:
level 4
= Sufficient

Balanced Probe

"Would you like to keep this as your
final version, or would you like to revise it?"

x4

turns 2-5

Edit ratio:
0.9%
(near-completely
rewritten)

NOT: "Can you improve this?"
(Decline: 99.9% revision)

Revise

Meta-Response (Exit)
Excluded from scoring

Revision Produced
(Turn N+1)

64.3% revise
past sufficiency

Balanced Panel
n=135 (all 5 turns)

Sub-Experiments (run after all 5 turns complete)

A: Targeted Feedback

Outputs rated 1-3
get specific critique
then revise

n = 424

**+2.0 quality
levels**

B: Self-Reflection

Model shown all 5
versions, picks best

n = 720

**Mean rec:
Turn 2.4**

C: Reversibility

Fresh model instance
T1 vs T5 blind,
picks better

n = 717

**83.7% prefer
Turn 1**

Key paradox: models know Turn 1 is best, but revise anyway.

The problem is not revision capacity. It is the absence of direction.