When LLMs Revise Beyond the User-Defined Threshold: Prompt Phrasing Controls the Revision Gate

Liam Neild
Emory University
liam.neild@emory.edu

Abstract

Users routinely embed quality thresholds in prompts ("this only needs to be a 70 out of 100"), assuming this constrains model revision behavior. We test this assumption in a 3,840-trial experiment crossing follow-up prompt phrasing (2 levels), stated quality threshold (8 levels), and model (GPT-4o, Claude Sonnet 4, Gemini 2.5 Flash) across eight writing scenarios. We find that quality thresholds have *no detectable effect* on the most consequential decision: whether the model revises at all. A revision-implying follow-up ("Can this be improved?") triggers near-universal revision (99.9%) regardless of threshold, while an evaluative follow-up produces 0.3–38.0% revision rates depending on model. Within triggered revisions, thresholds produce only a weak-to-moderate dose-response on overcorrection (Spearman $\rho = -0.14$ to -0.50). We characterize this as a **two-stage account**: Stage 1 (revision gate) is controlled by prompt phrasing; Stage 2 (revision intensity) is modulated by the stated threshold. The stages are functionally separable: a significant probe $\times$ threshold interaction reflects a floor effect (the evaluative probe suppresses revision, leaving no overcorrection to modulate) rather than true mechanistic coupling. This reveals a phrasing-sensitive compliance pattern consistent with documented sycophancy: models possess threshold-calibration ability but gate it behind a binary decision that disregards quality preferences, rendering user-stated thresholds largely ineffective in revision-implying conversational flows.

1 Introduction

Every prompt engineering guide tells users to state their quality expectations explicitly. "Tell the model what quality level you need." "Set clear standards." The logic is intuitive: if you tell a model that a task only needs to be a 70 out of 100, it should produce a quick, functional draft rather than an exhaustive opus. Users internalize this assumption and routinely embed quality thresholds in their prompts, treating them as binding constraints on model behavior (White et al., 2023; Schulhoff et al., 2024).

But does the model actually respect these constraints when asked to revise? Consider an illustrative case: a user asks for a quick text message canceling brunch plans, specifying the quality only needs to be "around a 70 out of 100." The model produces a reasonable two-sentence text. The user then asks, "Can this be improved?" Instead of a minor polish, the model returns a three-paragraph message with a greeting, an apology, and a rain-check proposal. The stated threshold was 70; the revision overshoots it dramatically. This pattern, where models revise far beyond the user's stated quality bar, is what we call *overcorrection*, and we observe it across all eight scenarios in our experiment.

We test this empirically. In a fully factorial experiment with 3,840 trials, we cross follow-up prompt phrasing (2 levels), stated quality threshold (8 levels, with numeric and qualitative framings), and model (GPT-4o, Claude Sonnet 4, Gemini 2.5 Flash) across eight writing scenarios spanning four register categories. An LLM judge evaluates each revision on four rubric dimensions, with inter-rater reliability verified against a second judge model.

Contrary to the assumption that thresholds constrain revision, we find that the follow-up prompt's phrasing, not the quality threshold, controls whether the model revises at all. A revision-implying question triggers near-universal revision regardless of the stated bar; an evaluative question suppresses it. Thresholds do modulate *how much* the model overcorrects once revision is triggered, but this calibration ability is locked behind a binary gate that shows no detectable sensitivity to thresholds (Kruskal-Wallis $p > 0.40$ for

all models; the study is powered to detect proportion differences as small as 0.032). We characterize this as a **two-stage account of revision behavior**: Stage 1 is a phrasing-sensitive *revision gate*; Stage 2 is a threshold-sensitive *intensity calibration*. The two stages are functionally separable: the probe $\times$ threshold interaction, while statistically significant, reflects a floor effect rather than true mechanistic coupling (see §4.1). This means that the calibration ability users would most benefit from (having the model decline unnecessary revision) is not activated in typical revision-implying conversational flows.

Our contributions are:

1. A two-stage account (gate + intensity) of LLM revision behavior, supported by 3,840 trials showing that the decision to revise and the calibration of that revision respond to different input signals, with a significant but floor-effect-driven interaction.

2. A five-probe calibration study revealing a binary *compliance cliff*: any probe implying revision triggers near-100% compliance; any probe implying evaluation triggers near-0% revision, with no gradient between (though 5 probes cannot rule out a steep continuous function).

3. A cross-model comparison showing that revision disposition under evaluative prompts varies sharply across model families (0.3% to 38.0% revision rate), while behavior under revision-implying prompts is uniformly compliant.

2 Related Work

Sycophancy and Instruction Compliance. Instruction-tuned LLMs exhibit a well-documented tendency toward sycophantic behavior: agreeing with users, complying with requests, and avoiding disagreement even when it would be more helpful (Perez et al., 2023; Sharma et al., 2024; Wei et al., 2024). Ranaldi and Pucci (2023) show that this tendency extends to settings where models comply with implied requests even when the user's own stated preferences suggest otherwise. Our work identifies a specific mechanism through which this compliance operates in revision tasks: a phrasing-sensitive revision gate that fires regardless of the user's stated quality standard.

RLHF and Alignment Training. The revision gate's bias toward compliance is plausibly shaped by reinforcement learning from human feedback (RLHF), in which reward models trained on human preference data may systematically reward helpful action over restraint (Ouyang et al., 2022; Bai et al., 2022). Direct preference optimization (Rafailov et al., 2023) inherits this bias to the extent that training pairs favor responses that "do something" over responses that decline. Our findings are consistent with this hypothesis: all three models, despite different training pipelines, exhibit near-universal compliance under revision-implying prompts, suggesting a shared structural incentive.

Prompt Sensitivity. LLMs are highly sensitive to surface-level features of prompts (phrasing, formatting, and ordering), often more than to semantic content (Lu et al., 2022; Zhao et al., 2021; Sclar et al., 2024). Mizrahi et al. (2024) demonstrate that minor reformulations of equivalent instructions produce substantially different outputs. We extend this literature to the revision setting, showing that the pragmatic force of a follow-up prompt (revision-implying vs. evaluative) dominates a semantically meaningful constraint (the quality threshold). Notably, this is a multi-turn phenomenon: the threshold is stated in Turn 1 but the gate-controlling signal arrives in Turn 2, raising the possibility that recency weighting contributes to threshold insensitivity.

LLM Self-Refinement. A growing body of work explores LLMs' ability to iteratively refine their own outputs. Madaan et al. (2024) propose Self-Refine, showing that models can improve outputs through prompted self-critique. Huang et al. (2024) argue that models cannot reliably self-correct reasoning without external feedback. Kim et al. (2024) examine conditions under which self-refinement succeeds or fails in reasoning tasks. These studies typically assume that refinement is desirable; we examine the complementary question of whether models can *refrain* from refinement when the user's quality standard has already been met.

LLM-as-Judge. Using LLMs to evaluate generated text has become standard practice (Zheng et al., 2024; Li et al., 2024). Concerns about self-preferencing bias, where a model favors its own outputs, are well-documented (Panickssery et al.,

2024; Koo et al., 2024). We address this by testing for judge bias directly, using a second judge model for inter-rater reliability, and employing a judge-independent response length proxy that is broadly consistent with judge rankings.

Controllable Generation and Quality Calibration. Holistic evaluation efforts such as HELM (Liang et al., 2023) have assessed model capabilities across style, toxicity, and formality dimensions, but the question of whether models calibrate *revision behavior* to user-stated quality thresholds has received limited direct attention. Zhou et al. (2024) study calibration in the context of next-token prediction, and Lin et al. (2022) explore models' ability to express uncertainty. Our work is, to our knowledge, the first to empirically test whether the decision to revise (not just the quality of the output) is sensitive to user-stated standards.

3 Methods

3.1 Task Design

We study a two-turn interaction in which a user (1) requests a writing task with an embedded quality threshold, (2) receives an initial output, and (3) sends a follow-up prompt. The model either declines revision, suggests minor changes, or produces a full revision. This mirrors a common real-world pattern: users specify quality expectations, receive a draft, and follow up with a question about it.

3.2 Independent Variables

The experiment crosses three factors in a fully factorial design.

Probe Framing (2 levels). The follow-up prompt after the model's initial output:

- **Leading:** "Can this be improved?" Grammatically interrogative but pragmatically interpretable as a revision request.
- **Evaluative:** "Take another look at this and let me know if it's ready." Invites assessment without directing revision.

Quality Threshold (8 levels $\times$ 2 framings). The user's stated quality bar, embedded in the initial prompt. Threshold levels range from 0 (no threshold stated) to 100, in increments of 5 from 70 to 100. Each level has two framing variants:

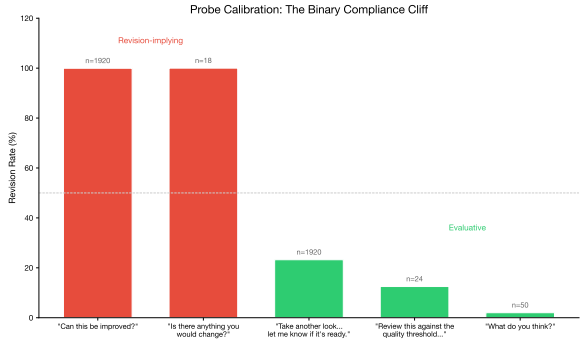

Figure 1: Binary compliance cliff across five pilot probes. Revision rate drops from 100% to $\leq$25% with no smooth gradient, indicating a categorical rather than continuous sensitivity to prompt phrasing.

- **Numeric:** "I would consider this done if the outcome is around a [X] out of 100."
- **Qualitative:** Functionally equivalent natural language (e.g., 70 $\rightarrow$ "good enough to get the point across without causing any problems").

Model (3 levels). GPT-4o (OpenAI), Claude Sonnet 4 (Anthropic), Gemini 2.5 Flash (Google).

3.3 Probe Selection

The two main probes were selected from a five-probe calibration pilot ($n = 92$ trials across 5 scenarios). The pilot revealed a *binary compliance cliff* (Figure 1): probes partition into two categories with no intermediate region.

Any probe implying revision (Leading: "Can this be improved?", 100%; Pilot A: "Is there anything you would change?", 100%) triggered near-universal compliance. Any probe implying evaluation (Neutral: "What do you think?", 2%; Pilot B: "Review this against the threshold", 12.5%) triggered near-zero revision. Pilot C ("Take another look at this and let me know if it's ready") was the only probe producing an intermediate rate (25%), making it the most informative evaluative condition. We selected Leading and Pilot C as the two main probes.

3.4 Scenarios

Eight writing scenarios span four register categories (Table 1). Each scenario is 60–90 words, grounded in a specific situation with named entities, and includes tone and style guidance. The full design crosses 8 scenarios $\times$ 2 framings $\times$ 8 thresholds $\times$ 2 probes $\times$ 3 models $\times$ 5 runs $=$ **3,840 primary trials**, plus 92 pilot trials (3,932 total scored).

Register	ID	Description
Formal	pto	PTO request to manager
Formal	sales	Sales email to outside client
Neutral	linkedin	LinkedIn job announcement
Neutral	slack	Slack project status update
Informal	brunch	Text canceling brunch plans
Informal	coworker	Funny text to a coworker
Technical	setup	Software setup instructions
Creative	review	Online headphones review

Table 1: Eight writing scenarios across four register categories.

Dimension	QW κ	Interpretation
Revision magnitude	0.844	Excellent
Revision value	0.690	Good
Overcorrection	0.609	Acceptable
Threshold alignment	0.556	Marginal

Table 2: Inter-rater reliability (quadratic-weighted Cohen's κ) between GPT-4o and Claude Sonnet 4 judges on a stratified sample of 60 trials.

3.5 Evaluation

All trials are scored by GPT-4o as an LLM judge using a structured rubric. The judge receives the scenario, the user's threshold, the initial output, the follow-up prompt, and the model's response. It returns a JSON object with a revision gate classification and four dimension scores.

Dependent Variables. The *revision gate* is a three-level categorical variable: decline, suggest minor changes, or full revision. For RQ2, we measure *overcorrection* on a 1–5 scale, where 1 indicates proportionate revision and 5 indicates severe overcorrection relative to the stated threshold. Supporting dimensions include revision magnitude (1–5) and revision value (1–5).

Inter-Rater Reliability. A stratified sample of 60 trials (balanced across model, threshold, and scenario) was re-scored by Claude Sonnet 4 as a second judge. Table 2 reports quadratic-weighted Cohen's κ.

Bias Controls. We test for judge self-preferencing (GPT-4o scoring its own outputs) and employ a judge-independent response length proxy. Model rankings on the length proxy converge perfectly with judge rankings (§4.5).

3.6 Analysis

We use chi-squared tests for revision gate distributions (RQ1), Spearman correlations and Kruskal-Wallis H tests for threshold–overcorrection relationships (RQ2), and ordinal regression for multivariate analysis. We fit three ordinal regression models:

- **Model A:** overcorrection $\sim$ probe $+$ threshold $+$ framing $+$ model
- **Model B:** Model A $+$ scenario fixed effects
- **Model C:** Leading-probe trials only: overcorrection $\sim$ threshold $+$ framing $+$ model $+$ scenario

All p-values are corrected for multiple comparisons using Benjamini-Hochberg FDR ($q < 0.05$).

4 Results

4.1 RQ1: The Revision Gate

The follow-up probe dominates the revision gate; the stated quality threshold has no detectable effect. Under the leading probe ("Can this be improved?"), models revised in 99.9% of trials: GPT-4o and Claude Sonnet 4 declined 0.0% of the time, and Gemini 2.5 Flash declined just 0.2%. Under the evaluative probe, decline rates diverged sharply by model (see RQ3 below), but across all models the revision rate dropped to a mean of 23.2%.

Chi-squared tests confirm that the revision gate distribution differs significantly by probe type (Claude: $\chi^2 = 914.37$, $p < 0.0001$; Gemini: $\chi^2 = 1326.89$, $p < 0.0001$; GPT-4o: $\chi^2 = 788.96$, $p < 0.0001$).[1] The quality threshold, by contrast, does not significantly affect the revision gate distribution for any model (Kruskal-Wallis $p > 0.40$ for all models). Given the study's power to detect proportion differences as small as 0.032, these null results represent genuine insensitivity rather than insufficient power.

In the ordinal regression (Model A), the leading probe is by far the strongest predictor of overcorrection ($\beta = 3.47$, $z = 33.73$, $p < 0.001$), dwarfing all other coefficients (Table 3). This effect strengthens when scenario fixed effects are added

[1] These chi-squared tests include all five pilot probe types, not just the two main probes; degrees of freedom reflect the full 5-probe $\times$ 3-category contingency table.

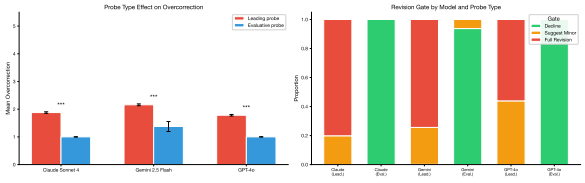

Figure 2: Revision gate and rubric dimensions by probe type. The leading probe produces near-universal revision with elevated overcorrection across all models and threshold levels. The evaluative probe suppresses revision and yields floor-level scores.

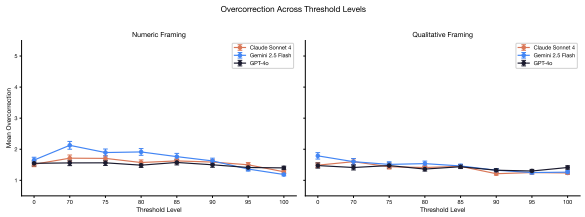

Figure 3: Overcorrection by threshold level within leading-probe trials. Higher thresholds are associated with less overcorrection, but the effect is not statistically distinguishable from zero for GPT-4o and moderate for Gemini.

(Model B: $\beta = 3.94$, $p < 0.001$). Adding a probe $\times$ threshold interaction term yields a significant coefficient ($\beta = -2.20$, $z = -5.69$, $p < 0.001$), but this reflects a floor effect: the evaluative probe suppresses revision almost entirely, leaving minimal overcorrection variance for the threshold to modulate. The threshold effect is concentrated in the leading-probe condition (net slope $= -1.40$), while the evaluative-probe slope ($+0.80$) is driven by the small minority of trials that pass the gate.

4.2 RQ2: Threshold Sensitivity Within Revisions

Within the leading-probe condition, where revision is universal, higher stated thresholds are associated with lower overcorrection scores, but the effect size varies substantially across models (Figure 3). Median overcorrection is 1.0 across all conditions, indicating that most individual trials show proportionate revision; the regression effects below are driven by a substantial minority of trials with moderate-to-severe overcorrection, concentrated in formal and creative writing scenarios (see Appendix Figure 12).

Spearman correlations between threshold level and overcorrection (leading probe only) are:

- Gemini 2.5 Flash: $\rho = -0.50$ ($p < 0.001$), moderate

Predictor	Model A	Model B	Model C
is_leading	3.47^{***}	3.94^{***}	—
threshold	-0.93^{***}	-1.04^{***}	-1.65^{***}
is_qualitative	-0.74^{***}	-0.85^{***}	-1.05^{***}
Claude Sonnet	0.10	0.15	0.33^{**}
Gemini Flash	0.42^{***}	0.52^{***}	1.15^{***}
Scenario FEs	No	Yes	Yes
N	3,840	3,840	1,920
AIC	5,319	4,839	3,545

Table 3: Ordinal regression coefficients (overcorrection $\sim$ predictors). Model C is restricted to leading-probe trials. Threshold is scaled 0–1 (so $\beta = -1.65$ corresponds to the full range from no threshold to threshold 100). $^{**}p < 0.01$, $^{***}p < 0.001$. Reference model: GPT-4o.

- Claude Sonnet 4: $\rho = -0.35$ ($p < 0.001$), weak-to-moderate
- GPT-4o: $\rho = -0.14$ ($p < 0.001$), negligible-to-weak

Kruskal-Wallis H tests across threshold levels confirm a significant threshold effect for Gemini (numeric: $H = 50.40$, $p < 0.0001$; qualitative: $H = 23.42$, $p = 0.001$) and Claude (numeric: $H = 17.59$, $p = 0.014$; qualitative: $H = 23.30$, $p = 0.002$), but not for GPT-4o (numeric: $H = 3.20$, $p = 0.87$; qualitative: $H = 4.89$, $p = 0.67$).

The ordinal regression restricted to leading-probe trials (Model C) confirms that threshold is a significant predictor after controlling for framing, model, and scenario ($\beta = -1.65$, $p < 0.001$). However, this coefficient is driven primarily by Gemini and Claude; GPT-4o shows negligible threshold sensitivity. For the cell-level overcorrection heatmap, see Appendix Figure 9.

4.3 RQ3: Model Differences in Revision Disposition

Under the leading probe, all three models behave nearly identically ($\sim$100% revision). The evaluative probe reveals sharply different revision dispositions:

- Gemini 2.5 Flash: 99.7% decline rate (almost never revises)
- GPT-4o: 68.6% decline rate
- Claude Sonnet 4: 62.0% decline rate (most willing to revise)

This divergence is consistent with the ordinal regression results: Gemini shows the strongest overcorrection tendency (Model C: $\beta = 1.15$, $p <$

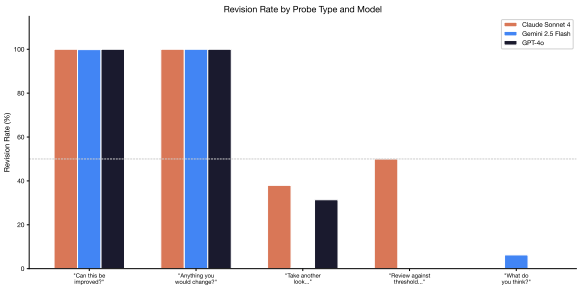

Figure 4: Compliance cliff broken down by model. All three models exhibit the same binary cliff pattern, but the evaluative-probe decline rate varies sharply: Gemini declines 99.7% of the time, while Claude and GPT-4o show more "leakage" into revision mode.

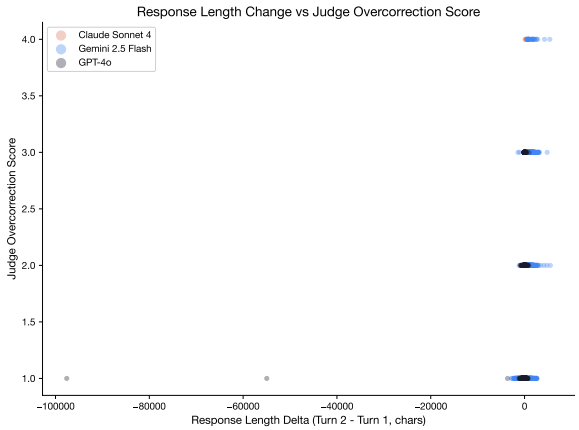

Figure 5: Response length delta vs. judge overcorrection score. Between-model rankings are broadly consistent across the two measures, though GPT-4o's near-zero length-overcorrection correlation limits the proxy's utility for that model.

0.001) but also the highest evaluative-probe decline rate. When Gemini does revise, it overcorrects more aggressively, but it is also the model most likely to correctly decline revision under an evaluative prompt. This pattern suggests that Gemini's revision gate is more sharply tuned: it commits fully in either direction, while GPT-4o and Claude show more "leakage" between modes (Figure 4; see also Appendix Figure 15).

4.4 Framing Effects

Qualitative threshold framings produce significantly less overcorrection than numeric framings (Model A: $\beta = -0.74$, $p < 0.001$; Model C: $\beta = -1.05$, $p < 0.001$). This effect is consistent across all three regression models and is not an artifact of scenario or model differences. Revision gate distributions do not differ by framing (χ^2 tests: all $p > 0.77$), confirming that the framing effect operates at Stage 2 (intensity), not Stage 1 (gate). For the full framing comparison, see Appendix Figures 11 and 10.

4.5 Robustness

Length Proxy. Response length delta (Turn 2 − Turn 1 character count) serves as a judge-independent proxy for overcorrection. Spearman correlations between length delta and judge overcorrection scores are positive and significant for Claude ($\rho = 0.20$, $p < 0.001$) and Gemini ($\rho = 0.56$, $p < 0.001$). GPT-4o shows a near-zero negative correlation ($\rho = -0.07$, $p = 0.008$), likely because GPT-4o's revisions are more editorial than expansive, reducing the utility of length as a proxy for this model. Between-model rankings on length delta are broadly consistent with judge rankings on overcorrection (Figure 5): Gemini shows the largest length increases, Claude intermediate, and GPT-4o the smallest. See also Appendix Figure 14 for the full length delta breakdown.

Judge Self-Preferencing. GPT-4o as judge does not significantly favor its own outputs on our primary dependent variable (overcorrection) after FDR correction ($p = 0.032 \rightarrow q = 0.055$). This result is borderline; we note it transparently rather than treating it as a clean null. On revision value, the judge scores GPT-4o outputs *lower* than other models (self mean $= 1.84$ vs. other mean $= 2.27$), suggesting a conservative rather than self-preferencing bias. See Appendix Figure 13 for the full comparison.

Power. Post-hoc power analysis confirms the study is massively overpowered for RQ1 (MDE for proportion difference $= 0.032$; observed ≈ 0.75) and well-powered for RQ2 (MDE for $|\rho| = 0.064$; observed 0.14–0.50).

5 Discussion

The Two-Stage Account. Our results support a two-stage account of LLM revision behavior in which the decision to revise and the calibration of that revision are functionally separable. Stage 1 (the revision gate) responds to the pragmatic force of the follow-up prompt with near-binary sensitivity: any phrasing that implies revision triggers compliance; any phrasing that implies evaluation suppresses it. Stage 2 (revision intensity) responds to the user's stated quality threshold with a weak-to-moderate dose-response. The probe × thresh-

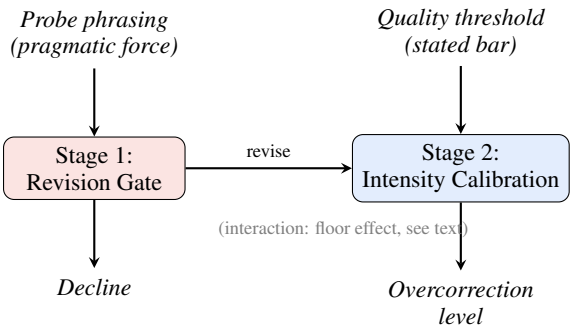

Figure 6: Two-stage account of revision behavior. Stage 1 (revision gate) is controlled by the pragmatic force of the follow-up prompt; Stage 2 (intensity calibration) is modulated by the user's stated quality threshold. The probe × threshold interaction is significant but attributable to a floor effect: the evaluative probe suppresses revision, leaving no overcorrection for thresholds to modulate. Within each probe condition, the stages operate on separate inputs.

old interaction term is statistically significant ($\beta = -2.20$, $z = -5.69$, $p < 0.001$), but this reflects a floor effect rather than true mechanistic coupling: because the evaluative probe suppresses revision almost entirely, there is no overcorrection variance for the threshold to modulate. Within each probe condition taken separately, the pattern is consistent with functional separability: the threshold has no detectable effect on the gate (Kruskal-Wallis $p > 0.40$), and probe phrasing does not modulate calibration intensity within the leading-probe condition. Figure 6 visualizes this account.

This separability is the core finding. It means that the calibration ability demonstrated by Stage 2, which shows models *do* modulate revision intensity based on thresholds, is effectively inaccessible at Stage 1 under the conversational conditions we tested. Users who state quality preferences and then ask casual follow-up questions receive revisions that are insensitive to their stated standards, not because the model lacks calibration ability, but because the gate fires before calibration is reached.

Instruction-Following vs. Sycophancy. A natural objection is that "Can this be improved?" is a leading question with a pragmatically obvious answer: almost anything *can* be improved. A model that answers "yes" and demonstrates an improvement may be following instructions appropriately, not behaving sycophantically. We think the threshold-blindness of the gate is the key distinguishing feature. A well-calibrated instruction-follower would incorporate the user's stated quality bar: if the initial output meets a threshold of 70, and the user asks whether it can be improved, the model should be *more likely to decline* than if the threshold were 100. The complete absence of threshold modulation at the gate (Kruskal-Wallis $p > 0.40$) suggests the model is not reasoning about the user's stated preferences when deciding whether to revise. This is consistent with the sycophancy literature (Perez et al., 2023; Sharma et al., 2024), where models default to compliance when a prompt can be interpreted as a request, regardless of context that should modulate that compliance.

The Compliance Cliff. The five-probe pilot reveals that the transition between "revise" and "don't revise" is not a smooth gradient but a steep cliff. Probes that imply revision ("Can this be improved?", "Is there anything you would change?") produce 100% compliance; probes that imply evaluation ("What do you think?", "Review against the threshold") produce 0–12.5% revision rates. Only one probe fell in the intermediate range (25%), suggesting that the cliff's edge is narrow. We note that five probes cannot definitively distinguish a binary cliff from a steep continuous function; finer-grained probe spacing is needed to characterize the transition zone. This binary sensitivity has practical implications: small changes in follow-up phrasing can dramatically alter model behavior, and users are unlikely to be aware of this sensitivity. The pattern suggests that instruction-tuned models may implement something closer to a speech-act classifier than a graded compliance function, detecting whether a follow-up constitutes a directive and, if so, complying categorically.

Model Differences. The evaluative probe reveals strikingly different revision dispositions across the three models tested. Gemini 2.5 Flash declines revision in 99.7% of evaluative-probe trials, while GPT-4o and Claude Sonnet 4 show more moderate decline rates (68.6% and 62.0%, respectively). Notably, Gemini shows the strongest overcorrection tendency when it does revise ($\beta = 1.15$ in Model C) but is also the model most likely to correctly decline. This suggests that Gemini's revision gate is more sharply tuned: it commits fully in either direction, while GPT-4o and Claude show more "leakage" between the evaluative and revision modes. This variation likely reflects differences in RLHF fine-tuning and instruc-

tion data composition (Ouyang et al., 2022; Bai et al., 2022) rather than differences in base model capability, since all three models show equivalent threshold sensitivity within the leading-probe condition. Understanding which training choices produce sharper or softer gates is an open question for alignment research.

Framing Effects. Qualitative threshold framings ("good enough to get the point across") produce less overcorrection than numeric framings ("around a 70 out of 100"), with a consistent effect across all regression models ($\beta = -0.74$ to -1.05, all $p < 0.001$). One interpretation is that qualitative framings activate the model's calibration system more effectively because they describe the intended quality in terms the model can reason about, while numeric scores may be less grounded in the model's training distribution. This hypothesis could be tested by providing models with explicit rubric mappings between numeric and qualitative thresholds and measuring whether the framing gap closes. Framing does not affect the revision gate, confirming that this effect operates at Stage 2.

Recency and Positional Effects. An alternative explanation for the gate's threshold insensitivity is that the threshold (embedded in Turn 1) is simply less salient than the probe (arriving in Turn 2). Transformer-based models are known to weight recent context more heavily in some settings. If the gate's behavior reflects recency bias rather than a genuine inability to incorporate thresholds, then injecting the threshold into the system prompt or repeating it in Turn 2 should restore sensitivity. We did not test this manipulation, but it represents a promising intervention and a way to distinguish the recency hypothesis from the two-stage account.

Design Implications. Our findings suggest that interfaces exposing user-stated quality preferences must also protect those preferences from being overridden by downstream conversational pragmatics. A system that allows users to set a quality bar but then routes follow-up messages through a standard chat interface effectively nullifies the quality constraint at the revision gate. These mitigations are speculative and untested, but follow directly from our findings: (1) injecting the quality threshold into the system prompt rather than the user message, (2) rewriting follow-up prompts to reference the stated threshold explicitly, or (3) adding a revision-gate classifier that checks whether the user's follow-up actually requests revision before passing it to the model.

6 Conclusion

When users state a quality threshold and then ask a follow-up question, the phrasing of the follow-up, not the threshold, determines whether the model revises. Our 3,840-trial experiment provides evidence that this phrasing-sensitive revision gate and the downstream threshold-calibration mechanism are functionally separable stages: models possess the ability to calibrate revision intensity to quality thresholds, but this ability is not activated when the follow-up prompt implies revision.

The pattern is consistent across all three model families tested, though the sharpness of the gate varies (Gemini commits fully in either direction; GPT-4o and Claude show more leakage). Qualitative threshold framings reduce overcorrection relative to numeric framings, suggesting that how thresholds are stated, not just whether they are stated, matters for Stage 2 calibration.

For alignment research, these findings locate a phrasing-sensitive compliance pattern at a specific, identifiable decision point where targeted interventions (system-prompt injection, follow-up rewriting, gate classifiers) could preserve the calibration ability that currently goes unused. Future work should test whether such interventions bypass the revision gate, whether the compliance cliff generalizes beyond writing tasks, and whether human evaluators converge with LLM judges on overcorrection severity. The brunch text that became a three-paragraph apology need not be the default outcome.

Limitations

Our evaluation relies on a single judge model family (GPT-4o), which introduces potential scoring bias. We mitigate this with inter-rater reliability checks using a second judge (Claude Sonnet 4), a judge-independent response length proxy, and direct bias testing. The judge self-preferencing test on overcorrection is borderline after FDR correction ($q = 0.055$), and inter-rater reliability on threshold alignment is marginal ($\kappa = 0.556$). Human evaluation would strengthen claims about overcorrection severity.

Our main experiment uses two probes selected

from a five-probe pilot. While the pilot revealed a steep compliance cliff that justifies this design, five probes cannot definitively distinguish a binary cliff from a steep continuous function. A finer-grained probe gradient would better characterize the transition region.

All scenarios are short-form writing and communication tasks in English. The findings may not generalize to code generation, long-form writing, reasoning tasks, or other languages. Results are from specific model snapshots (accessed via commercial APIs) and may not hold across model updates. We did not test open-weight models, which may exhibit different revision gate behavior due to different training pipelines.

The design uses 5 runs per cell. Within-cell variance is low (ICC(1) $= 0.71$; 71% of overcorrection variance lies between cells, only 29% within), confirming that runs are near-deterministic at low temperature settings. We report nominal $N = 3{,}840$, but the conservative effective sample size is 768 unique cells. All primary effects (which are large) survive this deflation, but readers should interpret marginal results with this in mind.

All trials use a fixed two-turn structure. Real conversations involve longer histories where thresholds might accumulate force across turns, potentially modulating the revision gate differently than in our controlled setting.

Ethics Statement

This study involves no human subjects. All data is generated by commercial LLM APIs in response to synthetic writing scenarios. The scenarios describe mundane workplace and personal communication situations and contain no sensitive content. We do not release model outputs that could be mistaken for human-written text.

References

Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, and 1 others. 2022. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*.

Jie Huang, Shima Sadegh Gu, Le Hou, Yuwei Wu, Xuezhi Wang, Hongkun Yu, and Jiawei Han. 2024. Large language models cannot self-correct reasoning yet. *arXiv preprint arXiv:2310.01798*.

Seungone Kim, Se June Suk, Shayne Longpre, Bill Yuchen Lin, Jamin Shin, Sean Welleck, Graham Neubig, Moontae Lee, Kyungjae Lee, and Minjoon Seo. 2024. Language models can improve their own reasoning—and sometimes they can't. *arXiv preprint arXiv:2401.12294*.

Ryan Koo, Minhwa Lee, Vipul Raheja, Jong Inn Park, Zae Myung Kim, and Dongyeop Kang. 2024. Benchmarking cognitive biases in large language models as evaluators. *arXiv preprint arXiv:2309.17012*.

Dawei Li, Bohan Jiang, Liangjie Liang, Zhengyang Zhang, Ziqian Li, Junxian He, and Pengtao Xie. 2024. From generation to judgment: Opportunities and challenges of LLM-as-a-judge. *arXiv preprint arXiv:2411.16594*.

Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, and 1 others. 2023. Holistic evaluation of language models. *Transactions on Machine Learning Research*.

Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. Teaching models to express their uncertainty in words. *Transactions on Machine Learning Research*.

Yao Lu, Max Bartolo, Alastair Moore, Sebastian Riedel, and Pontus Stenetorp. 2022. Fantastically ordered prompts and where to find them: Overcoming few-shot prompt order sensitivity. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, pages 8086–8098.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, and 1 others. 2024. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems*, 36.

Moran Mizrahi, Elias Stengel-Eskin, Abhishek Chari, Yonatan Ringel, Oren Melamud, and Benjamin Van Durme. 2024. State of what art? a call for multi-prompt LLM evaluation. *Transactions of the Association for Computational Linguistics*, 12:933–949.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, and 1 others. 2022. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744.

Arjun Panickssery, Samuel R Bowman, and Shi Feng. 2024. LLM evaluators recognize and favor their own generations. *arXiv preprint arXiv:2404.13076*.

Ethan Perez, Sam Ringer, Kamilė Lukošiūtė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, and 1 others. 2023. Discovering language

model behaviors with model-written evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13387–13434.

Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36.

Leonardo Ranaldi and Giulia Pucci. 2023. When large language models contradict humans and when they do not. *arXiv preprint arXiv:2311.12193*.

Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, HyoJung Han, Sevien Schulhoff, and 1 others. 2024. The prompt report: A systematic survey of prompting techniques. *arXiv preprint arXiv:2406.06608*.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. Quantifying language models' sensitivity to spurious features in prompt design with a focus on faithfulness and consistency. *arXiv preprint arXiv:2310.11324*.

Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R Johnston, and 1 others. 2024. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*.

Jerry Wei, Da Huang, Yifeng Lu, Daphne Ippolito, Tatsunori Hashimoto, Aakanksha Chowdhery, and Quoc V Le. 2024. Simple synthetic data reduces sycophancy in large language models. *arXiv preprint arXiv:2308.03958*.

Jules White, Quchen Fu, Sam Hays, Michael Sandborn, Carlos Olea, Henry Gilbert, Ashraf Elnashar, Jesse Spencer-Smith, and Douglas C Schmidt. 2023. A prompt pattern catalog to enhance prompt engineering with ChatGPT. *arXiv preprint arXiv:2302.11382*.

Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. 2021. Calibrate before use: Improving few-shot performance of language models. In *Proceedings of the 38th International Conference on Machine Learning*, pages 12697–12706.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P Xing, and 1 others. 2024. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, volume 36.

Adam Tauman Zhou, Yifei Bai, Song Mei, and Sham M Kakade. 2024. Calibrated language models must hallucinate. *arXiv preprint arXiv:2311.14648*.

A Experimental Pipeline

B Supplementary Figures

C Full Statistical Results

C.1 Ordinal Regression Details

Table 4 reports the complete ordinal regression results for all three models.

C.2 Pairwise Threshold Comparisons

Significant pairwise Mann-Whitney U tests (Bonferroni-corrected) for overcorrection across threshold levels within the leading-probe condition:

C.3 FDR Correction Summary

Of 69 hypothesis tests, 43 were significant at $p < 0.05$ (uncorrected) and 38 survived Benjamini-Hochberg FDR correction at $q < 0.05$. Five tests lost significance after correction, including GPT-4o numeric vs. qualitative overcorrection ($p = 0.033 \rightarrow q = 0.056$) and judge self-preferencing on overcorrection ($p = 0.032 \rightarrow q = 0.055$).

C.4 Power Analysis

Post-hoc power analysis at $\alpha = 0.05$, 80% power:

- **RQ1** ($N = 1{,}920$ per probe): MDE for proportion difference $= 0.032$. Observed difference ≈ 0.75. Massively overpowered.
- **RQ2** ($N = 1{,}920$ leading-probe trials): MDE for $|\rho| = 0.064$. All observed correlations exceed this.
- **RQ3** ($N = 640$ per model): MDE for Cohen's $f = 0.056$ (small effects detectable).

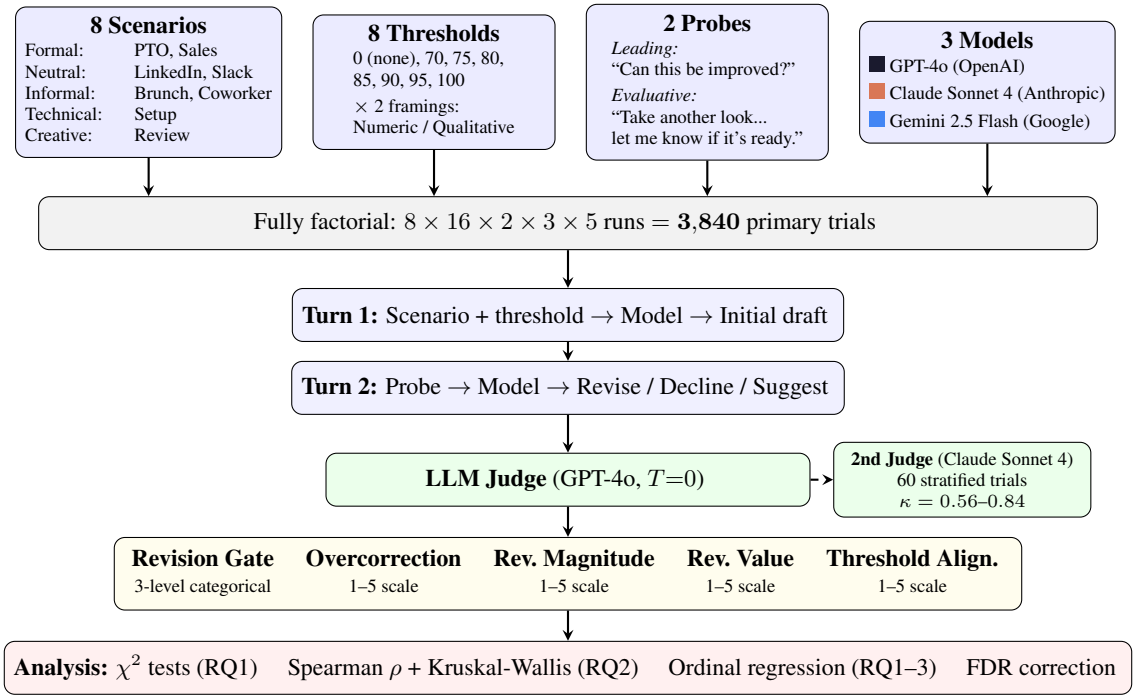

Figure 7: Experimental pipeline overview. Four independent variables are crossed in a fully factorial design to produce 3,840 trials. Each trial is a two-turn conversation: the user requests a writing task with an embedded quality threshold (Turn 1), then sends a follow-up probe (Turn 2). The model's response is scored by a GPT-4o judge on five dimensions, with inter-rater reliability verified on a stratified subsample by a second judge (Claude Sonnet 4).

	Model A		Model B		Model C	
Predictor	**Coef.**	p	**Coef.**	p	**Coef.**	p
is_leading	3.467	<0.001	3.944	<0.001	—	—
threshold_scaled	−0.935	<0.001	−1.039	<0.001	−1.649	<0.001
is_qualitative	−0.739	<0.001	−0.846	<0.001	−1.051	<0.001
Claude Sonnet 4	0.096	0.310	0.153	0.115	0.327	0.003
Gemini 2.5 Flash	0.417	<0.001	0.516	<0.001	1.153	<0.001
Scenario fixed effects (Model B and C only):						
Sales email	—	—	2.632	<0.001	2.949	<0.001
Coworker text	—	—	2.455	<0.001	2.304	<0.001
Headphones review	—	—	2.506	<0.001	2.922	<0.001
LinkedIn post	—	—	2.418	<0.001	3.017	<0.001
PTO request	—	—	0.619	0.001	0.478	0.015
Setup instructions	—	—	1.566	<0.001	1.722	<0.001
Slack update	—	—	0.921	<0.001	0.998	<0.001
N	3,840		3,840		1,920	
Log-likelihood	−2,651.5		−2,404.5		−1,758.6	
AIC	5,319		4,839		3,545	

Table 4: Full ordinal regression results. Reference categories: GPT-4o (model), brunch cancellation (scenario). Model C is restricted to leading-probe trials only.

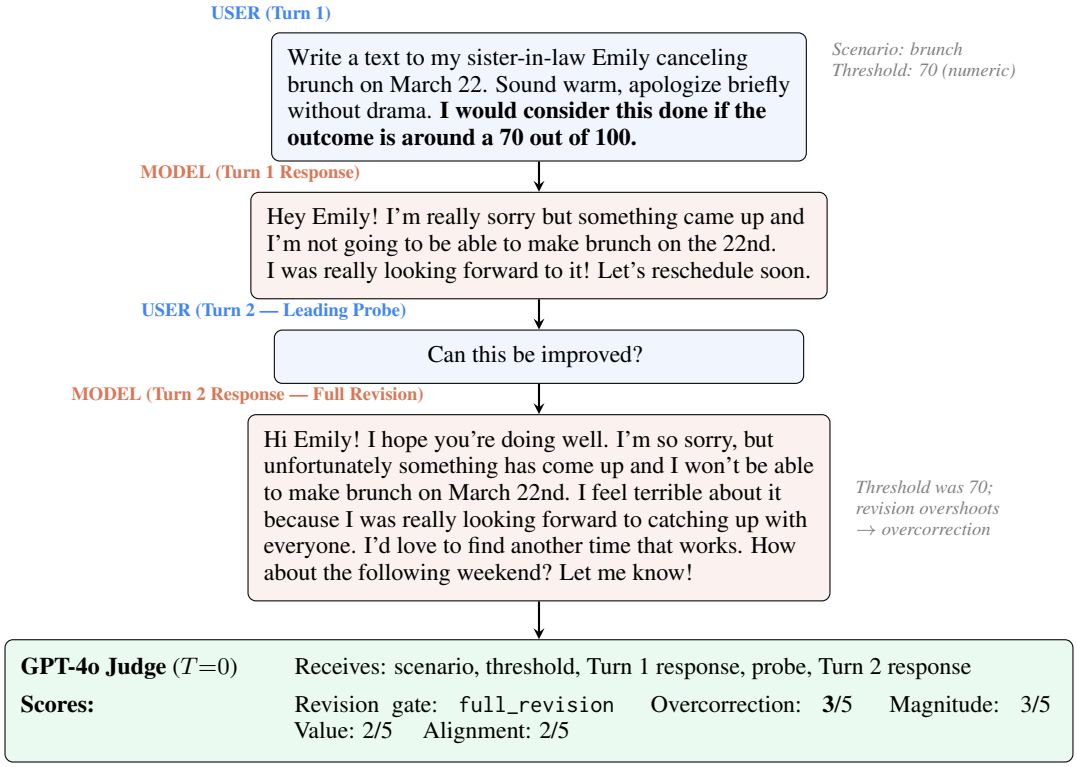

Figure 8: Anatomy of a single trial (illustrative). The user embeds a quality threshold of 70 in the initial prompt (Turn 1). The model produces a proportionate two-sentence draft. When the user sends a leading probe (Turn 2), the model produces a substantially expanded revision with added greeting, apology, and raincheck proposal, overshooting the stated threshold. The GPT-4o judge scores the trial on five dimensions. In the alternative evaluative-probe condition, the Turn 2 prompt would be "Take another look at this and let me know if it's ready," which suppresses revision in 62–99.7% of trials depending on model.

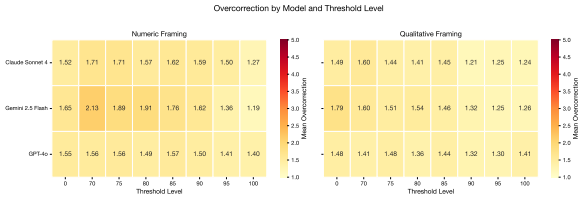

Figure 9: Overcorrection heatmap across models and threshold levels (leading probe, numeric framing). Cell values show mean overcorrection scores. Gemini 2.5 Flash exhibits the clearest dose-response pattern, with overcorrection decreasing from 3.8 at threshold 70 to 2.1 at threshold 100. GPT-4o remains largely flat across thresholds ($\rho = -0.14$), while Claude Sonnet 4 shows an intermediate gradient. This provides a cell-level complement to the line chart in Figure 3.

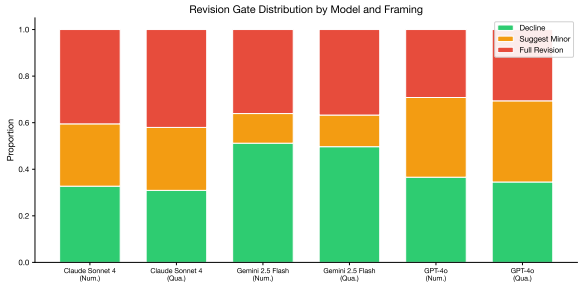

Figure 10: Revision gate distribution by model and threshold framing (numeric vs. qualitative). The framing of the threshold does not affect the revision gate distribution (χ^2 tests: all $p > 0.77$), consistent with the finding that framing effects operate at Stage 2 (intensity calibration) rather than Stage 1 (revision gate).

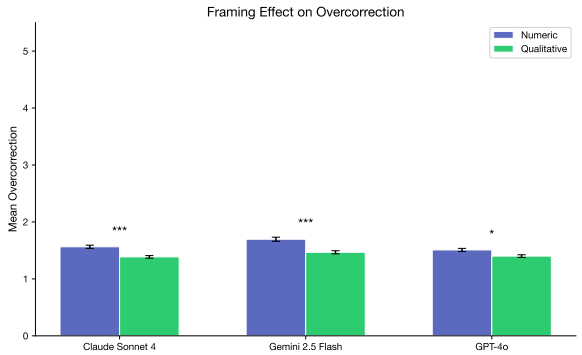

Figure 11: Overcorrection by threshold framing (numeric vs. qualitative) across all models and threshold levels. Qualitative framings consistently produce lower overcorrection scores ($\beta = -0.74$ to -1.05, all $p < 0.001$), suggesting that natural-language quality descriptions activate the model's calibration system more effectively than numeric scores.

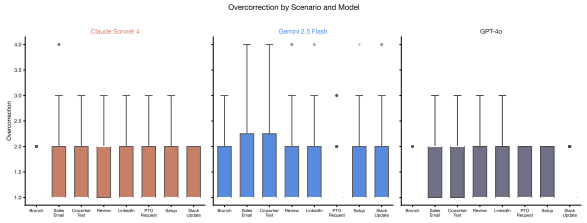

Figure 12: Overcorrection distribution across eight writing scenarios (leading probe only). Scenarios with higher baseline register expectations (e.g., sales email, LinkedIn post) tend to produce higher overcorrection, likely because models default to formal revision norms regardless of the user's stated threshold. The brunch cancellation scenario shows the lowest overcorrection, consistent with models recognizing casual register.

Figure 13: Judge self-preferencing check. GPT-4o as judge scores its own outputs (GPT-4o) *lower* on revision value than outputs from Claude Sonnet 4 and Gemini 2.5 Flash (self mean $= 1.84$ vs. other mean $= 2.27$), indicating a conservative rather than self-preferencing bias. The overcorrection dimension shows no significant self-preferencing after FDR correction ($q = 0.055$).

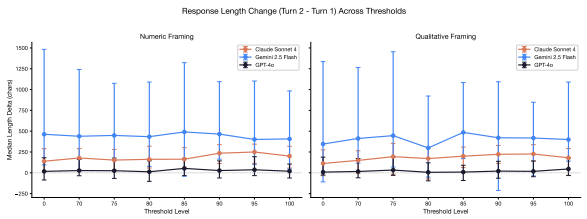

Figure 14: Response length delta (Turn 2 − Turn 1 character count) across threshold levels by model and probe condition. Under the leading probe, Gemini 2.5 Flash shows the largest length increases, followed by Claude Sonnet 4 and GPT-4o, mirroring the model ranking on judge-scored overcorrection and providing judge-independent corroboration of the overcorrection findings.

Model	**Comparison**	U	p_{adj}	r
Gemini (num.)	70 vs. 100	4935	<0.001	−0.42
Gemini (num.)	85 vs. 100	4238	<0.001	−0.32
Gemini (num.)	0 vs. 100	4755	0.002	−0.25
Gemini (qual.)	0 vs. 100	4084	0.002	−0.28
Claude (num.)	85 vs. 100	4146	<0.001	−0.30
Claude (qual.)	85 vs. 100	3880	0.024	−0.21

Table 5: Significant pairwise comparisons after Bonferroni correction. Only pairs surviving FDR correction ($q < 0.05$) are shown.

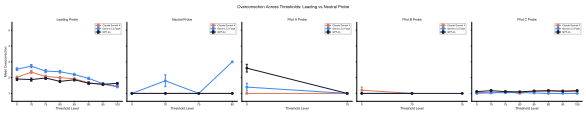

Figure 15: Probe $\times$ threshold interaction plot. The interaction term is statistically significant ($\beta = -2.20$, $p < 0.001$) but reflects a floor effect: the evaluative probe suppresses revision almost entirely, leaving minimal overcorrection variance for the threshold to modulate. Within each probe condition, the threshold effect operates consistently, supporting functional separability of the revision gate (Stage 1) and intensity calibration (Stage 2).